# GPU Sharing Mechanisms for Multi-Tenant Inference Serving

Group A: Search Execution and Evidence Report

MUHAMMAD BURHAN UN NADEEM (SE-24071)
MUHAMMAD FAAIZ ALAM (SE-24094)
ALI RAZA (SE-24302)

11 October 2026

Operating Systems Complex Engineering Problem
NED University of Engineering and Technology

**Abstract**

This intermediate deliverable records execution of the submitted SLR protocol and prepares traceable evidence for the final CEP report. Two databases yielded 689 captured query occurrences; supplementary public venue and author searches supplied additional named candidates. Twenty-one distinct peer-reviewed studies currently meet full-text eligibility. Unavailable full texts, incomplete publisher coverage and extraction gaps remain explicit. The paper set and conditional synthesis are provisional; the completed four-database SLR and final CEP submission are not claimed.

**Status: evidence checkpoint, not final SLR.** The publication window is 1 January 2020–11 October 2026 inclusive. One AI reviewer conducted screening and appraisal at the user's direction; the listed students are the project authors, not claimed independent reviewers.

# Contents

**1 Purpose and requirements** **3**

**2 Protocol adoption and deviations** **3**

**3 Search execution and provenance** **3**

**4 Screening checkpoint and reconciliation** **4**

**5 Current accepted primary studies** **5**
5.1 Source-anchored finding register . . . . . 7

**6 Review of six supplied full texts** **8**

**7 Quality appraisal and extraction controls** **9**

**8 Provisional answers to the research questions** **10**
8.1 RQ1: mechanism and control boundaries . . . . . 10
8.2 RQ2: outcome definitions and comparability . . . . . 10
8.3 RQ3: operating conditions . . . . . 10
8.4 RQ4: trade-offs and counterevidence . . . . . 10
8.5 RQ5: conditional engineering choices . . . . . 11

**9 Sensitivity, limitations and continuation** **11**

**10 Evidence package and institutional requests** **12**

# 1 Purpose and requirements

The engineering problem is to share a physical GPU among inference tenants while balancing useful work and density against predictable latency and isolation. The assigned alternatives are time-slicing, MPS, MIG and API remoting/virtualization, including hybrids. The unit of technical comparison is local sharing on one GPU. A study containing cluster experiments qualifies only when local inference-sharing evidence is identifiable; cluster-wide GPU counts are never substituted for local occupancy.

The submitted `report/SLR_protocol.tex` and `report/SLR_protocol.pdf` are the protocol baseline identified by the user. The protocol contents are preserved. Query deliverables reside in the root `report/` folder; Group A evidence remains under `groups/group-a/evidence/`. The final assignment still requires 15–20 unique peer-reviewed primary studies, at least 15, a 12-point IEEE-cited LaTeX report, and approximately 20–30 pages of main content. Meeting a provisional paper count does not complete searching or the final deliverable. No additional intermediate-page requirement was supplied.

The five questions ask which mechanisms are evaluated (RQ1), how outcomes are measured and compared (RQ2), which operating conditions change results (RQ3), what trade-offs and gaps remain (RQ4), and what conditional choice is supported (RQ5). All five map to extraction and the claim register, rather than to an unconditional ranking of mechanisms.

# 2 Protocol adoption and deviations

The Group A execution copy adopts the submitted protocol. Amendment GA-A001 extends the cutoff from 27 September to 11 October 2026. GA-A002 records the user's instruction that Codex conduct screening and quality appraisal as a single AI reviewer. There is no independent second reviewer, inter-rater agreement statistic or adjudicated human consensus. GA-A003 permits supplementary full-text work while unavailable database coverage remains open. GA-A004 records the actual public-interface execution and metadata capture in place of signed-in publisher exports. None changes the inclusion criteria to meet a target count.

Eligibility requires an English peer-reviewed original journal or full conference report, a local sharing/partition/process/API mechanism, explicit concurrent inference or justified mixed inference evidence, and sufficient configuration/outcome detail. Training-only, cluster-only, surveys and vendor claims are excluded from the primary set. Unretrieved reports are pending access outcomes, not assessed full-text rejections. Related conference, journal and preprint reports are one underlying study until an independent contribution is established.

# 3 Search execution and provenance

NED's E-Resources page lists the four planned publisher platforms. The listing was checked on 11 October 2026: `https://eaklibrary.neduet.edu.pk/eakl/E-Resources.html`. It confirms platform availability in the list, not entitlement to each requested article. No university credentials or licensed full-text PDFs enter this public repository.

| Source/run | Actual implementation | Observed hits |
|---|---|---|
| IEEE MPS | All metadata; years 2020–2026 | 25 |

| Source/run | Actual implementation | Observed hits |
|---|---|---|
| IEEE MIG | All metadata; years 2020–2026 | 23 |
| IEEE temporal | All metadata; years 2020–2026 | 45 |
| IEEE API | All metadata; years 2020–2026 | 9 |
| Springer master | General search; English; research articles/conference papers;2020–2026 | 587 |
| ACM pilot | Search blocked by challenge loop | Unknown; no import |
| ScienceDirect pilot | Search page unavailable | Unknown; no import |

The common block is GPU or graphics-processing-unit AND inference, model serving, deep learning, DNN or LLM. IEEE uses one family block per run. Springer uses all family blocks in one OR group. Exact strings, filters, source URLs, dates and export limitations are in `search_runs.csv`. The IEEE field implementation differs from the planned field-qualified Command Search and is disclosed, not described as an equivalent validated command. Research type/language and precise 2026 cutoff checks are manual where facets are unavailable. No accepted study is dated after 2025.

All 30 Springer result pages (20 hits/page,7 on the final page) and all IEEE pages were captured, yielding 587+25+23+45+9=689 occurrences. These are rendered-result metadata imports, not publisher-generated RIS/CSV exports. Springer export redirects to sign-in; IEEE results export requires personal sign-in and public citation export timed out. The imports retain IDs, database/run, page, title, DOI/URL and displayed date/item metadata. Clock time was not retained and is marked unrecorded.

Springer's broad index includes reference/full-text matches and unrelated acronym expansions such as welding MIG. Search hits therefore do not imply eligible GPU studies. Seed checks recovered ParvaGPU and KRISP in IEEE; papers in other publisher proceedings require supplementary discovery and completion of ACM coverage. Supplementary records are selected named-paper lookups checked after adoption, including submitted seeds and related systems. They are not an exhaustive search-engine result export or completed backward/forward citation search. Their follow-up query strings are reproducible suggestions, not invented historical exact queries.

# 4 Screening checkpoint and reconciliation

Every formal occurrence receives a record ID. Exact DOI/URL and normalized-title checks link duplicate occurrences without deleting provenance. Bibliographic report IDs identify candidate reports; study IDs identify accepted underlying studies. The REEF journal extension is selected for GA-S001; its less complete conference report is a linked E8 companion and adds no study.

| Checkpoint quantity | Count |
|---|---|
| Formal query occurrences | 689 |
| Duplicate occurrences | 10 |
| Unique formal records | 679 |
| Selected supplementary records | 20 |

| Checkpoint quantity | Count |
|---|---|
| Excluded before full text | 575 |
| Full texts assessed | 28 |
| Full-text exclusions | 7 |
| Accepted unique studies | 21 |
| Pending / not retrieved | 96 |

The reconciled captured subset is: 689-10+20=699 records; 699 equals pre-full-text exclusions plus assessed reports plus pending records. Assessed reports equal included reports plus full-text exclusions. Missing ACM/ScienceDirect populations are unknown and outside these arithmetic identities. This checkpoint is intentionally not a completed PRISMA2020 flow. Reports sought include pending retrieval attempts, but complete formal-database retrieval and final study-version resolution remain open.

Title decisions are first-pass relevance judgments, explicitly labelled as title stage. Abstract decisions cite publisher scope; full-text decisions identify a section or experiment. Training and inference are not interchangeable: PipeSwitch and SIRIUS are retained for explicit inference outcomes, whereas TGS is rejected for training-only evidence. The Fermilab paper was retained after its multi-model slice experiment was inspected, rather than merely because it mentioned MIG. DIPPM was rejected because profile prediction alone does not test concurrent tenant sharing.

Full-text rejection examples include DIPPM (no concurrent tenant comparison), the accelerator-pool vRAN simulation and the low-latency FEC offloading paper (no inference evidence), DRMaestro's Rodinia workload evaluation (no serving inference), CalcGraph (assigned GPU partitioning left to future work), and the mediated-device passthrough benchmark (HPC applications without inference). The complete reasons and stages are in `rejected.csv`. Pending files have no fabricated technical results or QA scores.

# 5 Current accepted primary studies

The following studies are accepted for this intermediate evidence set. Selection remains subject to access completion and checking more complete peer-reviewed versions. KRISP is a scope-boundary inclusion under protocol 7.2 because it evaluates local spatial resource control and inference workers; its emulated AMD feature and non-NVIDIA baseline are explicitly isolated in synthesis. No vendor document or preprint-only unpublished work counts toward this table.

| Study | Report/year | Mechanism / scope | QA |
|---|---|---|---|
| GA-S001 | REEF 2025 [1] | Reset-based temporal preemption + dynamic kernel padding; NVIDIA binary instrumentation | 8.5 |
| GA-S002 | GPULET 2022 [2] | MPS + temporal | 8.5 |
| GA-S003 | PIPESWITCH 2020 [3] | Temporal context switching | 9.0 |
| GA-S004 | CLOCKWORK 2020 [4] | Temporal exclusive scheduling | 8.5 |
| GA-S005 | TORPOR 2025 [5] | API remoting + temporal swapping | 8.5 |

| Study | Report/year | Mechanism / scope | QA |
|---|---|---|---|
| GA-S006 | STREAMBOX 2024 [6] | API/runtime sandbox virtualization | 8.0 |
| GA-S007 | KRYPTON 2025 [7] | Kernel API virtualization + hybrid | 9.0 |
| GA-S008 | SIRIUS 2025 [8] | Temporal + MPS memory handover | 8.5 |
| GA-S009 | TALLY 2025 [9] | CUDA API interception + priority | 9.0 |
| GA-S010 | PARVA 2024 [10] | MIG + MPS | 8.0 |
| GA-S011 | ORION 2024 [11] | CUDA API interception + temporal | 9.0 |
| GA-S012 | GSLICE 2020 [12] | MPS adaptive provisioning | 9.0 |
| GA-S013 | KRISP 2023 [13] | Spatial partition/API control (AMD) | 8.0 |
| GA-S014 | XSCHED 2025 [14] | Temporal API interposition | 9.0 |
| GA-S015 | PAELLA 2023 [15] | Temporal software scheduling + MPS comparison | 8.5 |
| GA-S016 | GNNTRITON 2024 [16] | MIG + concurrent model serving | 8.5 |
| GA-S017 | VASTPIPE 2025 [17] | MPS adaptive quotas; separate AMD CU-mask stratum | 6.5 |
| GA-S018 | ARKGPU 2023 [18] | CUDA API interception + NVML monitoring and admission throttling | 7.0 |
| GA-S019 | DYNAP 2025 [19] | MIG dynamic placement/reconfiguration with reconfiguration-last policy | 7.0 |
| GA-S020 | IGNITER 2023 [20] | MPS active-thread quotas + interference-aware batch/provisioning and shadow process | 8.5 |
| GA-S021 | INFERFAIR 2024 [21] | Temporal SLO-aware scheduling + feedback goodput balancing; MPS comparator | 7.5 |

USENIX proceedings PDFs are official published copies. ACM/IEEE author or camera-ready copies are linked to verified proceedings metadata and do not become separate preprint studies. ParvaGPU's author manuscript is provisional for publisher-version comparison; its reported runtime-version anomaly remains flagged. The supplied REEF journal conversion establishes a more complete implementation/evaluation; it is now the preferred report for GA-S001. The conference report is retained as a linked companion. Accessible copies of Orion, GSLICE, KRISP, Tally and Paella avoid unnecessary institutional download requests.

## 5.1 Source-anchored finding register

| Study | Bounded finding | Source location |
|---|---|---|
| GA-S001 [1] | GV100 PyTorch DISB-C: Orion p 99 normalized latency 1.17x vs Reef below 1.05x; normalized total throughput 0.97x vs 1.14x(Fig 29). TensorRT comparison trades throughput: mean 1.21x Triton vs 1.11x Reef(Fig 28). All ratios retain model-solo normalization; do not mix engines/AMD strata. | Article pp 1:23-1:24,1:31-1:34; sections 6.2,6.6,7; Figs 28-29; Markdown lines 1111-1180,1600-1820 |
| GA-S002 [2] | Local gpulets support interference-aware provisioning; cluster throughput is a different outcome. | PDF p 10 Table 2/3; section 5; local profiling section 3 |
| GA-S003 [3] | Fig 7: inference preemption latency reduced by up to 38.2ms. | PDF pp 11-13; sections 6-6.2; Fig 7 |
| GA-S004 [4] | Fig 5 measures goodput satisfying SLO, rather than raw request throughput. | PDF p 10; section 6.1; Fig 5 |
| GA-S005 [5] | BERT Table 4: native 42ms, remote 43ms, PCIe swap 144ms, NVLink swap 45ms. | PDF pp 10-11; Table 4; sections 7.2-7.3 |
| GA-S006 [6] | Table 2 evaluates runtime sandbox workflows on shared GPUs. | PDF p 10 onward; section 9; Table 2 |
| GA-S007 [7] | Inference-specific sharing comparisons; section 2.2 tests fault propagation. | PDF p 10 Tables 1-2; section 5; section 2.2 Fig 3 |
| GA-S008 [8] | Inference p 99/SLO outcomes accompany training progress; separate TTFT/TBT LLM case. | Section 6.1 Fig 10; section 6.6 |
| GA-S009 [9] | At 10% idle time: normalized throughput Tally 1.19 vs TGS1.65. | PDF pp 9-11; sections 5-5.3; Fig 6 |
| GA-S010 [10] | Fig 7 separates internal/external fragmentation; Fig 5 counts GPUs. | PDF p 8 onward; sectionIV; TableIV; Figs 5/7 |
| GA-S011 [11] | Interference-aware scheduling compares inference latency and BE throughput. | PDF pp 8-9; sections 5.3-6.1; Tables 1/3 |
| GA-S012 [12] | GUE measures efficiency; shadow instance loading takes about 12s in Fig 15. | PDF pp 8-12; section 4; Table 3; Fig 15 |
| GA-S013 [13] | Kernel-wise right-sizing evaluates throughput and p 95 inference latency. | PDF pp 9-10; sectionsV-VI; TableIII |
| GA-S014 [14] | Fig 15: relative p 99 1.53/1.51/1.07 respectively versus solo. | PDF pp 14-15; section 8 Case 3 Fig 15; section 9 |
| GA-S015 [15] | p 99 latency and throughput under bursty model mixes; Fig 12 includes MPS. | PDF pp 9-12; section 7 Table 3/Figs 11-12; section 8 |

| Study | Bounded finding | Source location |
|---|---|---|
| GA-S016 [16] | Fig 10: thrashing after two models/20GB or five/40GB. | HTML Multi-model Scaling/Fig 10; Triton implementation; limitations |
| GA-S017 [17] | A100 VGG16+GoogLeNet: throughput 2.7x no MPS and 1.5x default MPS; not a universal average or p 99 result. NLP/CV mean 2.43x no MPS, 1.45x default MPS. MI100 mean 1.89x no mask, 1.24x default mask. | Published pp 457-459; sections 6.7-6.9; Figs 14-16; uploaded Markdown lines 1150-1290 |
| GA-S018 [18] | Section 5.2 reports all 30 tested pairs satisfy the 2x-solo p 99 QoS target with ArkGPU. Low-contention pairs can lose about 5% QPS from throttling. Four-GPU app-mix goodput gains 94.98% vs k 8s-native and 38.65% vs MPS are separate system-level results. | Published pp 313-319; sections 4.3,5.1-5.4; Table 3; Figs 14-21; Markdown lines 920-1360 |
| GA-S019 [19] | Single-A30 batch experiments: smaller partitions increase concurrency but can increase mean JCT; not online p 99 evidence. Fig 8 broader job set reports 39.03% lower makespan vs StaCon,12.14% vs MISO and 14.7% higher normalized STP vs StaCon. Prediction-error sensitivity:15% perturbation adds 8.09% makespan. | PDF pp 11-16; sections 4.1-4.3.2; Table 2; Figs 8-13; Markdown lines 700-865 |
| GA-S020 [20] | Local four-model latency prediction error 1.53%-5.02%(Fig 13). Fig 14 deployment: iGniter 6 GPUs/0 SLO violations vs gpu-lets+8/3,FFD+5/10,GSLICE+6/3;25% lower hourly instance cost vs gpu-lets+, not 25% higher local utilization. Fig 17 shows an initial p 99 violation before shadow-process recovery. | Published pp 820-824; sections 4.2,5.1-5.4; Figs 13-18; Markdown lines 938-1245 |
| GA-S021 [21] | Local Table 3 load 3:1:2: unfairness 1.05 vs Clockwork 1.50; SLO goodput 524.29 vs 457.24req/s. Other load ratios give smaller differences. Local unfairness is max/min model slowdown(Eq 2), not p 99 or security. Four-GPU MAF outcome kept separate. | Published pp 17-18; sections 5.1-5.4; Table 3; Eq 2; Markdown lines 718-899 |

# 6 Review of six supplied full texts

Five reports add unique studies: VastPipe, ArkGPU, Dyna-P, iGniter and InferFair. REEF's journal report explicitly extends OSDI2022 with a complete NVIDIA implementation, additional kernel handling and newer comparisons; the journal is preferred within GA-S001. Six previously pending reports are now assessed; the earlier conference report becomes a full-text E8 companion. There are 21 accepted studies,28 assessed reports including 7 full-text exclusions, and 96 pending

access cases. Search occurrence and discovery counts are unchanged.

The source files are Markdown conversions. Two-column interleaving, missing plot pixels and damaged mathematical symbols prevent exact extraction of some equations, CDFs, error bars and configurations. Only readable labelled tables, captions and explicitly stated numeric results enter the register. Missing values are not interpolated. The receipt ledger records filenames and SHA-256 hashes; licensed source contents are not published.

Local GPU strata determine inclusion. VastPipe's NPU-cluster outcomes, ArkGPU's four-GPU goodput, Dyna-P's multi-GPU aggregates and cluster simulation, iGniter's fleet cost, and InferFair's four-GPU trace are distinguished from local results. ArkGPU qualifies through the explicit mixed inference/training exception; transfer is limited to priority inference protection. Dyna-P's job completion and makespan do not establish online p 99. InferFair's unfairness is a relative slowdown ratio and does not establish adversarial isolation.

Journal REEF's same-engine GV100 PyTorch comparison reports DISB-C p 99 below 1.05 times solo versus 1.17 for Orion, alongside separately normalized total throughput. Its TensorRT comparison also reports lower mean normalized throughput than Triton; preserving this counterexample prevents a universal superiority claim. iGniter's initial p 99 violation followed by shadow-process correction likewise qualifies any steady-state SLO statement. Source locations and QA rationales are recorded per report.

At the stricter QA threshold 8, VastPipe, ArkGPU, Dyna-P and InferFair are omitted; their mechanism-specific conclusions weaken. At the protocol low-quality threshold 6, none is omitted. This is a reporting-quality sensitivity, not a pooled causal result. The set now exceeds the protocol's 15–20-study target range; retain all eligible studies at this checkpoint and document any final narrowing transparently rather than removing evidence to fit a count.

# 7 Quality appraisal and extraction controls

Q1–Q10 cover tenant question, mechanism/configuration, environment, workload/arrivals, comparator fairness, metric clarity, measurement controls, result support, validity limits and reproducibility/interpretability. Scores are 1 adequate,0.5 partial and 0 absent; totals 8–10 are high,6–7.5 moderate and below 6 low. The single-reviewer scores, concise rationales and source anchors are in `quality.csv`. Scores assess reporting and interpretability; they do not prove unbiased results, system security, reproducible artifacts or complete searching.

Current totals range from 6.5 to 9.0. Four newly included studies score below 8; their configuration or measurement controls and missing plot data require caution. Environment and measurement-control fields are frequently partial. No authors' experiments or code artifacts were rerun. Artifact availability alone is not experimental replication. The extraction register captures source/version, mechanism, hardware/runtime, tenant workload, comparator, a bounded finding, caveat and precise source location. Fields not yet fully transcribed are labelled NOT EXTRACTED, distinct from a source's not-reported value. A further detailed extraction pass is needed before final numerical comparisons.

Isolation observations remain four different columns: performance interference, resource/memory separation, adversarial security and fault containment. Most current findings establish workload performance rather than comparative security or fault containment. The KRYPTON fault experiment is specifically identified; it does not validate all mechanisms or all

GPU generations.

# 8 Provisional answers to the research questions

## 8.1 RQ1: mechanism and control boundaries

The set contains temporal scheduling, MPS resource provisioning, MIG plus MPS, and API/runtime interception or remoting. Hybrids are counted as studies once, not once per mechanism. Software scheduling may control task admission, kernel dispatch, preemption or memory handover; these are different operating boundaries. A process-sharing prototype, a kernel-intercepted container and a hardware slice cannot be assigned the same security meaning. The detailed mapping is in `extraction.csv` and `claims.md`.

## 8.2 RQ2: outcome definitions and comparability

Keep requests/second, SLO-compliant goodput, allocated GPU count, GPU busy time, SM occupancy, memory utilization and paper-specific efficiency ratios separate. Likewise p 95, p 98, p 99, average latency, time-to-first-token and time-between-tokens are not substitute measurements. A reduction in context-switch latency does not by itself establish p 99 end-to-end latency isolation. Normalize only within an explicitly comparable experiment and retain the original denominator and statistic.

| Comparison check | Required condition before numerical comparison |
|---|---|
| Hardware | Same GPU/vendor/memory/profile or explicit indirect label; separate AMD emulation from NVIDIA mechanisms. |
| Workload | Same models/input sizes, batching, tenants and co-tenant mix; mixed training comparisons retain inference outcomes. |
| Load | Same open/closed-loop arrival process, rate, burstiness, warm-up and measurement window. |
| Comparator | Same resource budget; disclose reimplementations, wrappers and baseline adaptations. |
| Outcome | Same latency percentile/SLO and throughput denominator; occupancy and device-count reduction remain separate. |

## 8.3 RQ3: operating conditions

Current evidence suggests that model memory demand, launch granularity, arrival burstiness, batching, priority mix and transfer path alter the sharing result. GPU generation constrains which resource boundaries are available. Local slice experiments are informative for memory pressure, while cluster scaling results require their own placement/network interpretation. Recommendation confidence falls when driver/runtime, quota/profile and experiment windows have not yet been completely transcribed.

## 8.4 RQ4: trade-offs and counterevidence

Stricter priority protection may trade lower-priority throughput for predictability. Remoting or swapping adds path-dependent transfer cost. Resource partitioning may reduce packing

flexibility. Fine-grained control can require compiler changes, API interception, driver hooks or baseline adaptations. These implementation costs must accompany reported gains. The source-anchored counterexample and parameter-specific results reside in the extraction register, avoiding an invented common percentage improvement across studies.

### 8.5 RQ5: conditional engineering choices

The following are provisional reviewer interpretations, not validated deployment instructions.

| Operating objective | Candidate to investigate | Evidence / unresolved condition |
|---|---|---|
| Priority-sensitive inference | Temporal/API scheduling with admission and memory controls | GA-S001, S003, S004, S008, S009, S011, S014; direct per-paper tests, indirect transfer. |
| Density with controlled workload mix | Tuned MPS provisioning | GA-S002, 012; check memory/compute contention and reconfiguration cost. |
| Slice boundary with concurrent models | MIG, optionally MPS inside a slice | GA-S010, 016; profile, packing and memory pressure; no universal isolation result. |
| Serverless remoting and elasticity | API remoting/runtime virtualization | GA-S005, 006,007; runtime compatibility and transport budget. |
| Strong adversarial/fault guarantee | No universal selection supported | Comparative evidence insufficient; retrieve security candidates and distinguish failure types. |

## 9 Sensitivity, limitations and continuation

Removing QA<6 removes no current study, so that sensitivity test is uninformative. Removing provisional author-manuscript evidence weakens hybrid MIG and the recorded negative throughput comparison. Excluding AMD-only feature evidence limits cross-vendor conclusions. Removing all incompletely transcribed environment profiles would sharply reduce evidence; final recommendations require completing those fields. No pooled effect or universal winner is asserted.

The main limitations are incomplete ACM/ScienceDirect searches, single-AI-reviewer bias, inaccessible full texts, supplementary-selection bias, baseline and hardware heterogeneity, incomplete parameter extraction and unresolved report versions. The accessible set is concentrated in systems proceedings and older DNN workloads; LLM and adversarial isolation evidence requires further checking. A minimum paper count cannot offset these limitations.

Next, continue public-access recovery and available database exports, complete parameter extraction and rerun claims/sensitivity. The user supplied six of the ten requested papers; all six

are assessed. The remaining four are unavailable and no further paper download is requested. Link any new report to its underlying study before adding a citation. Exact 2026 dates must be verified. Scope changes still require a documented amendment and instructor guidance; this report does not relax eligibility.

# 10 Evidence package and institutional requests

The evidence directory supplies search runs, imported records, duplicates, all decisions and reasons, accepted/rejected/pending views, retrieval lookups, QA, extraction, claims and checkpoint counts. `references.bib` contains only the 21 currently chosen primary reports, each cited here. The spreadsheet is a review view; CSV/JSON records remain the authority. Scripts validate linkage and reconciliation and regenerate derived views.

The unresolved access backlog is in `download_requests.csv`; there are zero active requests. `download_requests.md` records the six received conversions and four unavailable original requests. No additional paper download is requested. All requested publisher platforms appear on NED's list, but article-level access must be checked using university access. PDFs should be supplied privately through the conversation, not committed to the public repository. Query exports must preserve exact syntax, collection, facets and counts; sending papers alone does not repair incomplete systematic search coverage.

# References


[1] M. Han, R. Chen, W. Shen, H. Zhang, J. Yang, and H. Chen, "Real-time, Work-conserving GPU Scheduling for Concurrent DNN Inference," *ACM Transactions on Computer Systems*, vol. 44, no. 1, pp. 1–42, 2025, art. no. 1; publication date November 2025 as printed in supplied report. [Online]. Available: https://doi.org/10.1145/3768622

[2] S. Choi, S. Lee, Y. Kim, J. Park, Y. Kwon, and J. Huh, "Serving Heterogeneous Machine Learning Models on Multi-GPU Servers with Spatio-Temporal Sharing," in *USENIX ATC*, 2022, pp. 199–216. [Online]. Available: https://www.usenix.org/conference/atc22/presentation/choi-seungbeom

[3] Z. Bai, Z. Zhang, Y. Zhu, and X. Jin, "PipeSwitch: Fast Pipelined Context Switching for Deep Learning Applications," in *USENIX OSDI*, 2020, pp. 499–514. [Online]. Available: https://www.usenix.org/conference/osdi20/presentation/bai

[4] A. Gujarati, R. Karimi, S. Alzayat, W. Hao, A. Kaufmann, Y. Vigfusson, and J. Mace, "Serving DNNs like Clockwork: Performance Predictability from the Bottom Up," in *USENIX OSDI*, 2020, pp. 443–462. [Online]. Available: https://www.usenix.org/conference/osdi20/presentation/gujarati

[5] M. Yu, A. Wang, D. Chen, H. Yu, X. Luo, Z. Li, W. Wang, R. Chen, D. Nie, H. Yang, and Y. Ding, "Torpor: GPU-Enabled Serverless Computing for Low-Latency, Resource-Efficient Inference," in *USENIX ATC*, 2025, pp. 597–612. [Online]. Available: https://www.usenix.org/conference/atc25/presentation/yu

[6] H. Wu, Y. Yu, J. Deng, S. Ibrahim, S. Wu, H. Fan, Z. Cheng, and H. Jin, "StreamBox: A Lightweight GPU SandBox for Serverless Inference Workflow," in *USENIX ATC*, 2024, pp. 59–73. [Online]. Available: https://www.usenix.org/conference/atc24/presentation/wu-hao

[7] S. Zhang, A. Xu, Q. Chen, H. Zhao, W. Cui, Z. Wang, Y. Li, L. Xiao, and M. Guo, "KRYPTON: Efficient Performance-Aware GPU Sharing with Compatibility and Isolation through Kernel Space Interception," in *USENIX ATC*, 2025, pp. 1003–1019. [Online]. Available: https://www.usenix.org/conference/atc25/presentation/zhang-shulai

[8] J. Wang, Y. Wang, M. Han, and R. Chen, "SIRIUS: Colocating ML Inference and Training with Fast GPU Memory Handover," in *USENIX ATC*, 2025, pp. 1657–1675. [Online]. Available: https://www.usenix.org/conference/atc25/presentation/wang-jiali

[9] W. Zhao, A. Jayarajan, and G. Pekhimenko, "Tally: Non-Intrusive Performance Isolation for Concurrent Deep Learning Workloads," in *ACM ASPLOS Volume 1*, 2025, pp. 1052–1068. [Online]. Available: https://doi.org/10.1145/3669940.3707282

[10] M. Lee, S. Seong, M. Kang, J. Lee, G.-J. Na, I.-G. Chun, D. S. Nikolopoulos, and C.-H. Hong, "ParvaGPU: Efficient Spatial GPU Sharing for Large-Scale DNN Inference in Cloud Environments," in *IEEE/ACM SC*, 2024. [Online]. Available: https://ieeexplore.ieee.org/document/10793202

[11] F. Strati, X. Ma, and A. Klimovic, "Orion: Interference-aware Fine-grained GPU Sharing for ML Applications," in *ACM EuroSys*, 2024, pp. 1075–1092. [Online]. Available: https://doi.org/10.1145/3627703.3629578

[12] A. Dhakal, S. G. Kulkarni, and K. K. Ramakrishnan, "GSLICE: Controlled Spatial Sharing of GPUs for a Scalable Inference Platform," in *ACM SoCC*, 2020, pp. 492–506. [Online]. Available: https://doi.org/10.1145/3419111.3421284

[13] M. Chow, A. Jahanshahi, and D. Wong, "KRISP: Enabling Kernel-wise RIght-sizing for Spatial Partitioned GPU Inference Servers," in *IEEE HPCA*, 2023, pp. 624–637. [Online]. Available: https://doi.org/10.1109/HPCA56546.2023.10071121

[14] W. Shen, M. Han, J. Liu, R. Chen, and H. Chen, "XSched: Preemptive Scheduling for Diverse XPUs," in *USENIX OSDI*, 2025, pp. 671–692. [Online]. Available: https://www.usenix.org/conference/osdi25/presentation/shen-weihang

[15] K. K. W. Ng, H. M. Demoulin, and V. Liu, "Paella: Low-latency Model Serving with Software-defined GPU Scheduling," in *ACM SOSP*, 2023, pp. 595–610. [Online]. Available: https://doi.org/10.1145/3600006.3613163

[16] C. Savard, N. Manganelli, B. Holzman, L. Gray, A. Perloff, K. Pedro, K. Stenson, and K. Ulmer, "Optimizing High-Throughput Inference on Graph Neural Networks at Shared Computing Facilities with the NVIDIA Triton Inference Server," *Computing and Software for Big Science*, vol. 8, 2024, art. no. 14. [Online]. Available: https://link.springer.com/article/10.1007/s41781-024-00123-2

[17] L.-X. Ma, L.-P. Wang, E. Shao, R.-Y. Cao, and G.-M. Tan, "VastPipe: A High-Throughput Inference System via Adaptive Space-Division Multiplexing for Diverse Accelerators," *Journal of Computer Science and Technology*, vol. 40, no. 2, pp. 444–463, 2025. [Online]. Available: https://doi.org/10.1007/s11390-024-3773-5

[18] J. Lou, Y. Sun, J. Zhang, H. Cao, Y. Zhang, and N. Sun, "ArkGPU: enabling applications' high-goodput co-location execution on multitasking GPUs," *CCF Transactions on High Performance Computing*, vol. 5, pp. 304–321, 2023. [Online]. Available: https://doi.org/10.1007/s42514-023-00154-y

[19] T. Adufu and Y. Kim, "Dyna-P: placement-aware dynamic partitioning for lightweight applications with modern GPUs," *Cluster Computing*, vol. 28, 2025, art. no. 608. [Online]. Available: https://doi.org/10.1007/s10586-025-05284-2

[20] F. Xu, J. Xu, J. Chen, L. Chen, R. Shang, Z. Zhou, and F. Liu, "iGniter: Interference-Aware GPU Resource Provisioning for Predictable DNN Inference in the Cloud," *IEEE Transactions on Parallel and Distributed Systems*, vol. 34, no. 3, pp. 812–827, 2023. [Online]. Available: https://doi.org/10.1109/TPDS.2022.3232715

[21] Y. Peng and H. Peng, "InferFair: Towards QoS-aware scheduling for performance isolation guarantee in heterogeneous model serving systems," *Future Generation Computer Systems*, vol. 150, pp. 10–20, 2024. [Online]. Available: https://www.sciencedirect.com/science/article/pii/S0167739X23003187